



Getting Over The Bar

Porting AMODE 31 Assembler Code to AMODE 64

Session 20937

Ray Mullins
Broadcom MSD
6 March 2024

- These techniques are presented as suggestions
- Every situation is different
- Some of these techniques may not be applicable to your situation
- YMMV
- Use at your own risk
- z/OS, z/VSE, z/VM, z/TPF, and IBM are registered trademarks of International Business Machines, Inc.
- VSEⁿ is a trademark of 21st Century Software, Inc.
- All other trademarks are properties of their respective owners

Background

- AMODE 64 provides opportunity for exponential expansion of data capacity
- Has been available for around two decades in z/OS and z/VM, and almost as long in z/VSE
- Several ISVs are implementing exploitation to provide a competitive advantage
- Less overhead than AR mode (no AX and LX)
 - But is compatible with dataspace code (can use access registers)

General notes

- This presentation is about converting assembler code, but also references interlanguage issues
- It is z/OS-centric
- Many of these techniques are valid for z/VSE + VSEⁿ, z/VM, and Linux on Z, and probably for z/TPF as well
- ISVs continue to port code to AMODE 64 when either being forced to (e.g., JES2 control blocks moving), or for constraint relief

Notes on my examples

- Glossary
 - HOB = high-order bit
 - LOB = low-order bit
 - ROT = rule of thumb
- I use relative branches; base-displacement branches work just as well
- I personally use SL(G)R for zeroing a register. You can use whatever you like:
SGR, XGR, LGHI, BRCTG Rn, *, BRCTH Rn, */SRAG Rn, Rn, 32 (this one might play havoc with your R4HA...)

64-bit machine instruction notes

- Majority of opcodes with a “G” in the middle or end are “Grande” instructions
- Target is 64 bits (whether register or storage)
- Examples
 - AGR (Add Grande Register) is a 64-bit version of AR
 - LGHI is a 64-bit version of LHI
 - STG stores 64 bits at the indicated address
- 64-bit storage opcodes are RXY or RSY format
 - $Y = \pm 2^{20}$ displacement ($-524,288 \leftrightarrow +524,287 / X'80000' \leftrightarrow X'7FFFF'$)
 - There are 16- and 32-bit analogous Y instructions (e.g., AHY)

64-bit machine instruction notes

- GF instructions: 64-bit target, 32-bit source
 - AGF[R] – Add Grande From Fullword [Register]
- LL– prefixed instructions: Load Logical (unsigned)
 - LLGF[R]: Load Logical Grande From Fullword [Register]: 32 bits in lower half, upper half set to zeros
 - LLGT[R]: Load Logical Grande Thirty-One Bits [Register]: 31 bits in lower half, bit 32 and upper half set to zeros
- LA[Y] operates on the lower half only in 24– and 31-bit AMODE
- B[AS]SM: to enter 64-bit mode, LOB must be on (odd address)

Conversion strategy from a previous employer

64-bit conversion strategy example

1. Try not to change any storage areas, widen pointers in control blocks, etc. unless absolutely required
2. Expand all internal stacking services to save/restore 64-bit registers
3. Implement BASSM/BSM for all inter-program linkage, calling services, etc.
4. Updateabend recovery to handle programs running in AMODE 64
5. Change all register-clearing code to 64-bit instruction equivalents
6. Clone executable macros to create 31-bit counterparts; change all programs to use clones; change back during conversion
7. Update 64-bit version of complex executable macros to work in AMODE 64
8. One by one, change outer routines to AMODE 64
9. One by one, change internal subroutines and control blocks to AMODE 64
10. **GOALS:** Implement virtual storage constraint relief by moving larger storage areas above the bar, makes RMODE 64 possible to implement

64-bit conversion general comments

- Small and large scope changes
- Perform one step at a time
- Full rebuild and regression test after steps that involve a global change (steps 1-6)
- Full regression tests after steps involving changing one program at a time (steps 7-10)

Steps 1, 2, 4–Addresses in control blocks

- Modified only if absolutely necessary, e.g.
 - Internal stacking architecture
 - Internalabend recovery control blocks
- Internal-use-only control blocks that require changes
 - Convert fullwords to doublewords
 - Watch for usage of HOB of addresses (look for J(N)M/J(N)P instructions)
 - Watch for alignment and slack bytes, especially if addresses have ALETs (creates doubleword/fullword combination)
 - May require some reorganization to reduce slack bytes
 - Misalignment no longer carries a performance penalty, but storage updates are not synchronized across CPU caches
 - Watch out for interlanguage requirements (especially pointer data types)

Step 1–Supporting API calls in AMODE 64

- External-facing control blocks for APIs
 - AMODE 24/31 API
 - Keep fullword addresses (don't change)
 - AMODE 64 or AMODE-agnostic API
 - Redefine fullword addresses as AD, where a 64-bit caller stores in the doubleword at the label, & a 31-bit caller stores in the fullword at the label
 - Agnostic requires testing AMODE in the entry logic
 - Recommend using at least a 144–byte save area, mapping to the proper format as required (read Chapter 2 of the *z/OS MVS Assembler Services Guide*, subtitled “Everything You Ever Wanted To Know About Save Areas, But Were Afraid To Ask”)

Step 1–Supporting API calls in AMODE 64

- API routine handling of different AMODEs
 - Two separate modules (maintenance headache)
 - One module with separate pathing (smaller headache)
 - AMODE-specific APIs assembled from one source
 - Use a modified CSECT name as the differentiator, e.g. -4 suffix
 - SETC variables under conditional assembly (&SYSAMODE and &SYSASCE), also good for AR mode
 - Based on CSECT name, OPSYN to G or GF
 - 31-bit, L → LGF
 - 64-bit, L → LG
 - Create macros for fullword instructions to substitute grande instructions

Step 2—Instruction Conversion

- Updating internal stacking is a good introduction to instruction conversion
- Binary arithmetic is usually straight forward
 - $X \rightarrow XGF$
 - $XR \rightarrow XGR$
- LOAD, however, is context sensitive
 - Discussed later

Step 3–BASSM/BSM

- Preparatory for AMODE 64
- For staying in AMODE 31, this requires HOB to be turned on in either ACON or VCON (binder option HOBSET), or dynamically set, otherwise code swaps to AMODE 24
- Update any macro used to wrap internal subroutine calls as well
 - This is a good initiation into the required instruction changes
- Perform full regression tests
 - Lingering AMODE 24 code might be affected

Step 3—RMODE 24 diversion

- Exploit RMODE(SPLIT) if you have 24-bit code, using a glue routine

- A DCB exit example (24 to 31)

DCBABXIT RSECT ,	Address stored in DCBE
DCBABXIT AMODE 24	Control in AMODE 24
DCBABXIT RMODE 24	Staying below the line
LARL R15,ABXIT31	Point to actual exit
SAM31 ,	Flip to AMODE 31
BR R15	Transfer in AMODE 31
ABXIT31 RSECT ,	
ABXIT31 AMODE 31	Execute in AMODE 31
ABXIT31 RMODE ANY	It can be anywhere*

- The binder creates an RICON for the LARL, and CSV properly sets its value upon loading
- LARL does not set either AMODE indicator
- Requires PDSE target (program object) for SYSLMOD

Step 4—Recovery

- Must deal with AMODE 64 abends
- Must deal with registers with information in the high half
- Testing
 - Have a way to trigger an abend (undocumented command, for example)
 - Add code to this method to abend in AMODE 64

Step 5-Converting zeroing registers

- This should be straight forward
- A low-execution-time-impact step
- Clearing of high halves should not affect your code, depending on use
 - Unless you have implemented your own version of the compilers' HGPR option
- Perform full regression tests!
 - Code not under your control could be affected if it does not properly manage the high halves

Step 6—Create 31-bit versions of macros

- Allows you to retain macro names for AMODE 64
- Copy macros to a new name to reflect AMODE 31
 - OURSRV31, for example
- Refactor source to use new name
- Build
- Deal with the ones that somehow got missed
- Build
- Test

Steps 7, 8, 9–AMODE 64 conversion

- Now the real fun begins!
- Analyze, analyze, analyze
- Work your way inward, heading towards common subroutines

First and foremost!

- `AMODE 64` for the C/RSECT, binder parameters/statements
 - Tells the binder to do things
 - CSV does things based on what the binder does
 - Otherwise, you must do it yourself
- `SYSSTATE AMODE64=YES`
 - For the macros. Think of the poor macros!

Temporary testing tip

- In entry code of outermost routine, add a LMH R0,R15,=16F'-1'
- Alternatively, for system-wide settings, use mostly undocumented but supported PARMLIB(DIAGxx) parameter TRAPS NAME(IeaInitRegsTask,IeaInitArSRB) (initializes both access and general-purpose-register high halves with X'FFFFFFFF' for TCBs and SRBs)
- Any subsequent AMODE 64 address reference that has not insured the purity of the high half willabend
- Arithmetic references – “unpredictable results may will occur”

Converting lower levels and calling

- Callers must use BASSM into 64-bit mode (but we did that already in step 3)
- LOB of address must be set on
- Fun things ensue if both low half HOB and LOB are on
 - You are now executing in 64-bit mode in the 2GB-4GB address range.
 - In z/OS, you are in the JRE work area with expected abends
 - In other IBM Z operating systems, you are in usable storage, although it may not be allocated

Modifying address constants

- ADCONs

- Internal usage

- $A(\text{FOO}) \rightarrow \text{AD}(\text{FOO})$
 - $31 \rightarrow 64$:
 $A(\text{FOO} + \text{X}'80000000') \rightarrow \text{AD}(\text{FOO} + \text{X}'1')$
 - $24|31 \rightarrow \text{Just } A \rightarrow \text{AD}$

- External service call

- Switch to equivalent AMODE 64 IBM services when provided

- VCONs

- Almost no worries! The Binder does it all for you (if option HOBSET is on, for AMODE 31 addresses)
 - Bit must be explicitly turned on for AMODE 64 (either with +1 in VCON or at runtime by OILL $R_n, \text{X}'0001'$)

AMODE 64 calling AMODE 31 caveat

- If calling system or other services via macros that do not support BASSM (e.g., SVC, BASR) that cannot be invoked in AMODE 64, bracket call with SAM31/SAM64
- Specify LINKINST=BASSM on CALL macro
- Default CALL parameter list generation is contingent on SYSSTATE settings
 - Not 64-bit or PLIST4=YES: 4 bytes, VL honored
 - 64-bit or PLIST8=YES: 8 bytes, VL not honored
 - ASC=AR is tied to the z/OS version...
 - Before 2.2, ALETs always follow addresses
 - Starting with 2.2:
 - Not 64-bit or PLIST4=YES, include ALETs unless PLISTARALETs=NO
 - 64-bit or PLIST8=YES, suppress ALETs unless PLIST8ARALETs=YES
- Make sure parameters passed are in below-the-bar storage

IBM (Assembler) Services

- Only a few authorized assembler services support AMODE 64 invocation at this time
- No DFSMS macro support
- Some services (Callable Services, LE, z/OS UNIX®) have 64-bit equivalents (usually with a "4" somewhere in the name)
- Very few services RMODE 64 invocation, but services are being updated with every release
 - If you feel a system service should support invocation from RMODE 64, open a case (direct statement made to me from IBM z/OS developers)

System services returned values

- IBM explicitly documents that all system services that return information will return in the lower 32 bits only, unless otherwise documented (e.g. IARV64 addresses)
- Important for return codes! Use LTGFR to ensure that the high half of R15 stays clean
- One important documented exception: dynamic 24– and 31–bit storage allocation services return clean 64–bit addresses
 - There are some others, like the R0 value from LOAD
- ROT: Use the same standard in internal routines as well as APIs

Loading addresses

- LLGT is usually your friend for dealing with fullword 31-bit addresses, but if the high-order-bit (HOB) could be on...
 - If it is merely an end-of-parameter list flag, LLGT
 - If it is an AMODE indicator, use LLGF, **not LGF! (Fun ensues!)**
 - If it is a flag, use LLGF, test with TMLH, then NILH X'7FFF' as soon as possible
- Sometimes LLGTR or LLGFR might be a better replacement for LR than LGR
- Context and following code are important!

Return code branch tables

- After BASSM/BSM, the return address is odd!
- Subtract 1 from the displacement in the B instruction

```
• BASR    R14,R15
  B       4(R15,R14)
  J       BAD
  WTO     'It worked!'
```

becomes

```
BASSM R14,R15
B      4-1(R15,R14)
J      BAD
WTO    'It worked!'
```

- There are also implications for a vectored return based on R14
- *Personally, I am not a fan of either technique.*

Converting instructions

- General ROTs
 - Binary numbers, straight switch to grande, e.g., L → LG
 - Binary data where propagation of the sign bit would be an issue, use LOAD LOGICAL, e.g., SLR/ICM B '0011' → LLGH
 - Pure 31-bit address: LLGT
 - Register pair arithmetic and shifts → grande and single instructions, e.g., D → DSG(F)
 - No more games with testing condition codes after logical arithmetic

Some instruction optimizations

- Reduce number of instructions
 - Copy a register's value and immediately shift: replace with SLLG/K, which have source and target registers
- Replace storage references with immediate instructions
 - IF CH,R0,GE,=H'16' → IF CGHI R0,GE,16
 - The above is an example of the HLASMTK SPM. I highly recommend using them.

Instructions

- ICM[Y]
 - With B'1111'
 - If testing signed binary arithmetic value for (N)Z or (N)M, replace with LTGF
 - If testing unsigned binary arithmetic value for (N)Z, replace with LLGF/LTGR
 - If testing 31-bit address for (N)Z
 - If HOB might be on, LLGF/LTGFR
 - If HOB will never be on, LTGF
 - Unknown or don't care, LLGT/LTGFR
 - Unsigned B'0001' into a zeroed register, LLGC/LTGR
 - Signed B'1000'+SRA 24, LB/LTGR

Instructions

- ICM[Y] (continued)
 - Unsigned B'0011' into a zeroed register, LLGH/LTGR
 - B'0111' into a zeroed register
 - Don't change
 - B'1000' (MVCL padding value, sometimes X'80' into the HOB)
 - 0ILH/IILH
- IC[Y] into a zeroed register
 - LLGC

Instructions

- Add/Subtract instructions with constants for address manipulation
 - Consider converting to LA/LAY
 - Makes register usage more intuitive/self-documenting

Instructions—possible “gotchas”

- L (LOAD FULLWORD)
 - Context is important!
 - LGF, LLGF, LLGT are most frequent replacements
 - LGF if value is known to be a signed binary arithmetic value
 - LLGF if value is known to be unsigned and HOB is part of the value, or is a flag and needs to be tested first
 - LLGT if 31-bit address and HOB is irrelevant
 - Also applies to RR flavors of instructions, although by this point you probably have already handled them

Instructions—possible “gotchas”

- Halfword arithmetic
 - z14 *finally* introduced ?GH instructions!
 - Before that, only LGH/LLGH are available, so...
 - If you have control of the target field, consider expansion to fullword or doubleword
 - Remember step 1 and control blocks caveat
 - Otherwise, find an unused register, L(L)GH/?GR
- LM
 - IBM does not provide LMGF
 - If all are signed binary, LMH/SRAG 32
 - If all are unsigned, LM + LMH, or LMD, using =nF'0'
 - If mixed...split up

Instructions—possible “gotchas”

- BCTG and its relatives
 - 64-bit unsigned value
 - Exposure to horrendous CPU loop if high half polluted
 - Generally, avoid changing unless you can be absolutely sure the value will be within an acceptable range
 - HLASMTK DO SPM construct does not use it
- CKSM, Crypto, CU*, MVCLE and its relatives, PERFORM_RANDOM
 - Maximum lengths are governed by AMODE
 - Possible to process 8 exbibytes if you're not careful

Virtual Storage Constraint Relief-a goal!

- Storage can now be above the bar
 - Affects any control block that has pointers
- Modify control blocks first
 - a few related pointers at a time
 - Fully test
 - Lather, rinse, repeat
- Once done making pointer changes, begin changing dynamic storage allocation to 64-bit
 - Again, modify only one or related storage areas at a time
 - Test
 - Lather, rinse, repeat

Architecture level set issues

- As an ISV, you must consider how far back are your supported operating systems
 - Newest unsupported z/OS version 2.3 requires z12 or better; oldest supported version 2.4 also requires z12 or better
 - For z/VM, all supported releases require a z13 or better
 - For z/VSE/VSEⁿ, 6.1 and z10; 6.2 and 6.3, z196/114 (z11)
- z12 level set restrictions:
 - HLASM: OPTABLE(ZS6 or Z12)/MACHINE(ZSERIES-6)
 - z/OS XL C and Metal C, and other compilers: ARCH(10)
- z13 HLASM OPTABLE(ZS7 or Z13)/MACHINE(ZSERIES-7)

Architecture level set issues (con't)

- If you need an instruction that is not supported by an older architecture
 - Create macros to generate equivalent instructions
 - Test O' and define the macro if the opcode is not defined
 - When you change OPTABLE/MACHINE, the O' test becomes true and skips defining the macro.

z/VM, z/VSE, and z/TPF

- z/VM supports AMODE 64 very nicely
- z/VSE and VSEⁿ have limited support
 - Code can run AMODE 64 and obtain storage above the bar using IARV64 memory objects, but...
 - z/VSE and VSEⁿ system services do not support being called in AMODE 64
 - HLASM, compilers, LNKEDT do not support AMODE 64
 - Talk to 21st Century if you want that, I have a contact 😁
- z/TPF supports AMODE 64 very nicely

Interlanguage issues

- z/OS XL C/C++
 - Compiling under ILP32, all pointers are 4 bytes
 - Compiling under LP64, all pointers are 8 bytes, unless declared with `__ptr32`
 - If dealing with 8-byte pointers in a control block under ILP32, create a typedef that uses either a struct for allocating 4 bytes slack followed by a 4 byte pointer, or a union with a `long long` and a pointer (depending how you set up the field).
 - Watch out for pollution from the HGPR(NOPRESERVE) option, which allows the compiler to use the high halves of the registers
 - Python & Java use 64-bit C for their native interfaces

Interlanguage issues

- Enterprise COBOL for z/OS
 - Version 6.4 understands 64-bit pointers
 - For older compilers, 8-byte pointers can be split into a group level with a POINTER FILLER and the actual POINTER
 - V5+: watch out for pollution from the HGPR(NOPRESERVE) option, which allows the compiler to use the high halves of the registers

Interlanguage issues

- Enterprise PL/I for z/OS
 - Supports AMODE 64 with version 5.1 and LP(64) option
 - POINTER(64), with POINTER(32) for 4-byte pointers
 - See *Programming Guide* for advice
 - For version 4.5
 - POINTER is 4 bytes
 - If dealing with 8-byte pointers in a control block, create a TYPE that uses a STRUCTURE for allocating 4 bytes slack and a 4 byte pointer
 - Watch out for pollution from the HGPR(NOPRESERVE) option, which allows the compiler to use the high halves of the registers

Interlanguage issues

- REXX (interpreter and compiler), CLIST (ugh)
 - Still living in AMODE 31
 - STORAGE() does understand 64-bit addresses
- VS FORTRAN—Really? YES! But...
 - Only lives in AMODE 31
 - Don't be confused by POINTER*8
 - First fullword is ALET (used in EMODE)
 - Second fullword is address
- APL2, PASCAL/VS, VS/BASIC, RPG II, TSO CLIST
 - Just checking if you're still awake

Interlanguage issues

- z/OS Language Environment (including z/VM)
 - Does support AMODE 64
 - Provides AMODE 64 services
 - Well-documented
- z/VSE and VSEⁿ languages, including REXX (not assembler)
 - Still AMODE 31

Wrapping up

- Changing AMODEs is not a short project, by any means
- Test, test, test!
- After completion, (almost) no more limits
- It could be a competitive advantage to handle exponentially more data
- Did I mention test, test, test?

Questions? Comments? Blank stares?

For more information

- The up-to-date version of this presentation can be found at You can email questions at mfasmcpgmr@gmail.com
- You can find me on
 - LinkedIn: <https://www.linkedin.com/in/raymullins/>
 - Twitter (currently incorrectly known as X): <https://twitter.com/zarchasmpgmr>
 - Facebook Messenger, WhatsApp, Signal, Telegram, Keybase

Your feedback is important!

Submit a session evaluation for each session you attend!

This is session 20937

www.share.org/evaluation

SHARE mobile app

